

1 Theory Question

1. Conceptually speaking, how do we differentiate between the inliers and the outliers when using RANSAC for solving the homography estimation problem using the interest points extracted from two different photos of the same scene?

Answer: The RANSAC algorithm involves taking N trials to estimate the best set of inlier points gathered from a randomly collected sample of n points. This is accomplished by computing a homography using the randomly collected sample of points. This homography will then be used to check if each correspondence lies within an acceptable range δ from the actual coordinate. If the correspondence lies within the acceptable range, it is labelled as an inlier. Otherwise, it is labelled as an outlier. With N trials, the algorithm will find enough inliers that will help compute the closest estimate of the true pairwise homography where the inliers are the true correspondences.

2. As you will see in Lecture 13, the Gradient-Descent (GD) is a reliable method for minimizing a cost function, but it can be excruciatingly slow. At the other extreme, we have the much faster Gauss-Newton (GN) method but it can be numerically unstable. Explain in your own words how the Levenberg-Marquardt (LM) algorithm combines the best of GD and GN to give us a method that is reliably fast and numerically stable at the same time

Answer: The Gradient Descent (GD) method is the optimal method when the estimation obtained from the linear least squares method is farthest from the actual local minimum of the convex function. The Gauss-Newton method is the quickest way to the local minimum but it is not always along the direction of the steepest descent thus making it unstable. However, it is an optimal method to use when the estimation is close to the local minimum. The Levenberg-Marquardt algorithm combines the best of both methods by adding a heuristic damping coefficient μ which when manipulated can steer between the Gradient Descent approach or the Gauss Newton approach. At the beginning, μ is set to an extremely large value so that the steps taken along the convex function mimics Gradient Descent but when the estimation gets closer to the local minimum, μ decreases thus allowing for the steps to mimic Gauss-Newton.

2 Programming Section

2.1 Theoretical Background

2.1.1 RANSAC

As explained in the theoretical question, The RANSAC algorithm involves taking N trials to estimate the best set of inlier points gathered from a randomly collected sample of n points. This

is accomplished by computing a homography using the randomly collected sample of points. This homography will then be used to check if each correspondence lies within an acceptable range δ from the actual coordinate. If the correspondence lies within the acceptable range, it is labelled as an inlier. Otherwise, it is labelled as an outlier. With N trials, the algorithm will find enough inliers that will help compute the closest estimate of the true pairwise homography where the inliers are the true correspondences.

The parameters chosen for the RANSAC algorithm are as follows:

- $n = 5$ since the recommendation is for n to be between $4 < n < 10$
- $\epsilon = 0.85$ where ϵ is the probability that a randomly chosen (x, x') is an outlier
- $p = 0.99$ where p is the probability that at least one of the N trials will be free of outliers in the calculation of the homography matrix
- N trials is calculated as

$$N = \frac{\ln 1 - p}{\ln 1 - (1 - \epsilon)^n} \quad (1)$$

- n_{total} is the number of correspondences
- $M = n_{total}(1 - \epsilon)$ where M is minimum value for the size of the inlier set for it to be acceptable
- $\sigma = 2$ where σ is a constant in the noise-induced displacement of the true location of a pixel which is modelled by the Gaussian function

$$g(\Delta x, \Delta y) = \frac{1}{2\pi\sigma^2} e^{-\frac{(\Delta x)^2 + (\Delta y)^2}{2\sigma^2}} \quad (2)$$

and σ is usually set between $0.5 < \sigma < 2$

- $\delta = 3\sigma$ which is commonly used

2.1.2 Least Square Estimation

Linear Least Square Estimation is used to estimate a starting homography consisting of all the inlier points. The homography H represents a linear mapping from one homogeneous coordinate to another in different planes of the same dimension. \vec{x}' is computed by multiplying H with \vec{x} . The matrix H can be expressed as

$$H = \begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & h_{33} \end{bmatrix}$$

Therefore, using the relation mentioned above, the homogeneous point in the range is given by:

$$x'_1 = h_{11}x_1 + h_{12}x_2 + h_{13}x_3$$

$$x'_2 = h_{21}x_1 + h_{22}x_2 + h_{23}x_3$$

$$x'_3 = h_{31}x_1 + h_{32}x_2 + h_{33}x_3$$

To obtain the coordinates in the range on the real plane, we have to divide the first two coordinates of the homogeneous coordinate by the third coordinate, thus giving

$$x' = \frac{x'_1}{x'_3} = \frac{h_{11}x_1 + h_{12}x_2 + h_{13}x_3}{h_{31}x_1 + h_{32}x_2 + h_{33}x_3}$$

$$y' = \frac{x'_2}{x'_3} = \frac{h_{21}x_1 + h_{22}x_2 + h_{23}x_3}{h_{31}x_1 + h_{32}x_2 + h_{33}x_3}$$

We can simplify the x' and y' coordinate to also include the actual coordinates in the domain on the real plane by dividing the RHS by x_3 as shown below

$$x' = \frac{h_{11}x + h_{12}y + h_{13}}{h_{31}x + h_{32}y + h_{33}}$$

$$y' = \frac{h_{21}x + h_{22}y + h_{23}}{h_{31}x + h_{32}y + h_{33}}$$

We make the above x' and y' equations into a system of linear equations:

$$h_{11}x + h_{12}y + h_{13} - h_{31}xx' - h_{32}yx' - h_{33}x' = 0$$

$$h_{21}x + h_{22}y + h_{23} - h_{31}xy' - h_{32}yy' - h_{33}y' = 0$$

Therefore, for each correspondence we get the following matrix equation to solve for H

$$\begin{bmatrix} x_1 & y_1 & 1 & 0 & 0 & 0 & -x_1x'_1 & -y_1x'_1 \\ 0 & 0 & 0 & x_1 & y_1 & 1 & -x_1y'_1 & -y_1y'_1 \end{bmatrix} \begin{bmatrix} h_{11} \\ h_{12} \\ h_{13} \\ h_{21} \\ h_{22} \\ h_{23} \\ h_{31} \\ h_{32} \end{bmatrix} = \begin{bmatrix} x'_1 \\ y'_1 \end{bmatrix}$$

In an over-determined system when we have more than four correspondences, the linear least square estimation can be calculated by using the pseudo-inverse

$$\vec{h} = (A^T A)^{-1} A^T b \quad (3)$$

2.1.3 Levenberg-Marquardt Algorithm

As already mentioned in the theory question, the Levenberg-Marquardt algorithm combines the best of Gradient Descent and Gauss-Newton by adding a heuristic damping coefficient μ which when manipulated can steer between the Gradient Descent approach or the Gauss Newton approach. At the beginning, μ is set to an extremely large value so that the steps taken along the convex function mimics Gradient Descent but when the estimation gets closer to the local minimum, μ decreases thus allowing for the steps to mimic Gauss-Newton.

For this homework assignment, the Levenberg-Marquardt algorithm is implemented by using the built-in SciPy function. The cost function is given by the error between the actual coordinates in the range space and the computed coordinates using the homography matrix H .

2.2 Source Code

```
1 # Name: Nikita Ravi
2 # Class: ECE 66100
3 # Homework #5
4 # Deadline: 09/28/2022
5
6 # Import Modules
7 import cv2
8 import math
9 import numpy as np
10 import os
11 from scipy.optimize import leastsquares
12
13 def getimages(path, task=1):
14     path = path + "/Task" + str(task)
15     images = [cv2.imread(path + "/" + img) for img in sorted(os.listdir(path))[:-1]]
16     return images
17
18 def displayimages(images, idx=None):
19     if(idx):
20         cv2.imshow("Window", images[idx])
21     else:
22         cv2.imshow("Window", images)
23     cv2.waitKey(0)
24     cv2.destroyAllWindows()
25     quit()
26
27 def computehomography(domaincoord, rangecoord, over=False):
28     # From homework #3
29     n = domaincoord.shape[1]
30     A = np.zeros((2*n, 8))
31     b = np.zeros((2*n, 1))
32     H = np.zeros((3,3))
33
34     for idx in range(n):
35         A[2*idx] = [domaincoord[0][idx], domaincoord[1][idx], 1, 0, 0, 0, (-domaincoord[0][idx] * rangecoord[0][idx]), (-domaincoord[1][idx] *
```

```

36     rangecoord[0][idx]])
37     A[2*idx + 1] = [0, 0, 0, domaincoord[0][idx], domaincoord[1][idx],
38     1, (-domaincoord[0][idx] * rangecoord[1][idx]), (-domaincoord[1][idx] *
39     rangecoord[1][idx])]
40     b[2*idx] = rangecoord[0][idx]
41     b[2*idx + 1] = rangecoord[1][idx]
42
43 h = np.array([])
44 if(over):
45     Aplus = np.linalg.pinv(np.matmul(A.T, A)) #pseudo-inverse of A
46     h = np.dot(np.matmul(Aplus, A.T), b)
47 else:
48     h = np.dot(np.linalg.pinv(A), b)
49
50 row = 0
51 for idx in range(0, len(h), 3):
52     spliced = h[idx:idx+3]
53     if(len(spliced) == 3):
54         H[row] = spliced.T
55     else:
56         H[row] = np.append(spliced, [1])
57     row += 1
58
59 return H
60
61 def establishcorrespondences(image1, descriptor1, keyPoint1, image2,
62 descriptor2, keyPoint2, path):
63     # Citation: https://docs.opencv.org/4.x/dc/dc3/tutorialpymatcher.html
64     bruteForceMatcher = cv2.BFMatcher() # It takes the descriptor of one
65     feature in first set and matches it with all other features in second set
66     using some distance calculation - the closest distance is returned
67     matches = bruteForceMatcher.match(descriptor1, descriptor2) # returns
68     matches where k is specified by the user
69     matches = sorted(matches, key=lambda x: x.distance)
70     domaincoord, rangecoord = [], []
71
72 for match in matches:
73     domainpoint = list(map(int, keyPoint1[match.queryIdx].pt))
74     domainpoint.append(1)
75     domaincoord.append(domainpoint)
76     rangepoint = list(map(int, keyPoint2[match.trainIdx].pt))
77     rangepoint.append(1)
78     rangecoord.append(rangepoint)
79
80 combined = np.hstack((image1, image2))
81 result = cv2.drawMatches(image1, keyPoint1, image2, keyPoint2, matches
82 [:100], combined, flags=cv2.DrawMatchesFlagsNOTDRAW_SINGLE_POINTS)
83 cv2.imwrite(path, result)
84
85 domaincoord, rangecoord = domaincoord, rangecoord
86 return domaincoord, rangecoord
87
88 def sift(image1, image2, idx1, idx2):
89     siftdetector = cv2.xfeatures2d.SIFTcreate() # Create SIFT Detector

```

```

82     keyPoint1, descriptor1 = siftdetector.detectAndCompute(image1, None)
83     keyPoint2, descriptor2 = siftdetector.detectAndCompute(image2, None)
84
85     return establishcorrespondences(image1, descriptor1, keyPoint1, image2,
86                                     descriptor2, keyPoint2, "hw05/siftcorrespondences/Task1/image" + str(idx1)
87                                     ) + "image" + str(idx2) + ".jpg")
88
89
90 def findinliers(H, domaincoord, rangecoord, delta):
91     rangehat = np.matmul(H, domaincoord.T).T
92     rangehat = rangehat.T / rangehat.T[2, :]
93     rangehat = rangehat.T
94     error = np.sum(np.abs(rangehat - rangecoord)**2, axis=1)
95     indices = np.where(error <= delta)[0]
96     return indices
97
98
99 def ransac(images, task=1, LM=True):
100    """
101    Constants
102    n = 5 (4 + n + 10)
103    ntotal = number of correspondences
104    epsilon = 0.85
105    p = 0.99
106    N = ln(1-p)/(ln[1-(1-epsilon) n])
107    M = (1-epsilon) * ntotal
108    delta = 3 * sigma - typical p.11.3
109    sigma = 2 - typical p11.3
110    """
111
112    # Constants
113    n = 5
114    epsilon = 0.85
115    p = 0.99
116    N = int(math.log(1-p) / (math.log(1-(1-epsilon)**n)))
117    sigma = 2
118    delta = 3 * sigma
119
120
121    # Inputs
122    images = [img.copy() for img in images]
123    gray = [cv2.cvtColor(img, cv2.COLORBGR2GRAY) if len(img.shape) == 3 else
124            img for img in images]
125    Hlist = []
126    bestinliers = []
127
128    for idx in range(len(gray) - 1):
129        domaingray, domainimage = gray[idx], images[idx]
130        rangegray, rangeimage = gray[idx + 1], images[idx + 1]
131        domaincoord, rangecoord = sift(domaingray, rangegray, idx, idx +
132                                         1)
133        domaincoord, rangecoord = np.array(domaincoord), np.array(
134                                         rangecoord)
135
136        ntotal = domaincoord.shape[0]
137        M = int((1 - epsilon) * ntotal)
138        inlierpercentage = 0
139        indices = list(range(ntotal))

```

```

131
132     bestinlieridx = None
133     for trial in range(N): # Conduct N Trials
134         randomsampleindices = np.random.choice(indices, n) # Randomly
135         select a sample of indices of n correspondences
136         randomdomains = domaincoord[randomsampleindices]
137         randomranges = rangecoord[randomsampleindices]
138         Hwithnoise = computehomography(randomdomains.T, randomranges.
139                                         T, over=False)
140         inlieridx = findinliers(Hwithnoise, domaincoord, rangecoord,
141                               delta)
142
143         numinlierindices = len(inlieridx)
144         if(numinlierindices / ntotal >= inlierpercentage):
145             inlierpercentage = numinlierindices / ntotal
146             bestinlieridx = inlieridx
147
148         if(len(bestinlieridx) < M):
149             print("WARNING: Number of inlier points is less than the minimum
150 requirement")
151
152         bestinliers.append(bestinlieridx)
153         H = computehomography(domaincoord[bestinlieridx].T, rangecoord[
154         bestinlieridx].T, over=True)
155         if(LM):
156             hvector = H.flatten()
157             result = leastsquares(costfunction, hvector, args=(domaincoord
158 [bestinlieridx], rangecoord[bestinlieridx]), method="lm")
159             hprime = result.x
160             H = hprime.reshape((3,3))
161             Hlist.append(H)
162             displayinliersoutliers(bestinliers, domaincoord, rangecoord,
163             domainimage, rangeimage, "hw05/inlieroutlier/Task" + str(task) + "/"
164             "image" + str(idx) + "image" + str(idx + 1))
165
166     return Hlist, bestinliers
167
168
169 def costfunction(h, domaincoord, rangecoord):
170     X = []
171     F = []
172     #[ (x,y,1), ...
173     for idx in range(len(domaincoord)):
174         xd = domaincoord[idx][0]
175         yd = domaincoord[idx][1]
176         xr = rangecoord[idx][0]
177         yr = rangecoord[idx][1]
178         X.append(xr)
179         X.append(yr)
180
181         F.append((h[0] * xd + h[1] * yd + h[2]) / (h[6] * xd + h[7] * yd + h
182 [8]))
183         F.append((h[3] * xd + h[4] * yd + h[5]) / (h[6] * xd + h[7] * yd + h
184 [8]))
185
186     return np.array(X) - np.array(F)

```

```

175
176 def displayinliersoutliers(inlierindices, domaincoord, rangecoord, image1
177   , image2, path):
178   combinedinliers, combinedoutliers = np.hstack((image1, image2)), np.
179   hstack((image1, image2))
180   height, width, channel = image1.shape
181   inlierindices = inlierindices[0].tolist()
182
183   for idx in range(domaincoord.shape[0]):
184     # Inliers
185     if(idx in inlierindices):
186       pt1 = tuple(domaincoord[idx])
187       pt2 = rangecoord[idx][0] + width, rangecoord[idx][1]
188
189       cv2.circle(combinedinliers, pt1[0:2], 3, (255, 0, 0), 1)
190       cv2.circle(combinedinliers, pt2[0:2], 3, (0, 255, 0), 1)
191       cv2.line(combinedinliers, pt1[0:2], pt2[0:2], (0, 0, 255), 1)
192
193     else:
194       # Outliers
195       pt1 = tuple(domaincoord[idx])
196       pt2 = (rangecoord[idx][0] + width, rangecoord[idx][1])
197
198       cv2.circle(combinedoutliers, pt1[0:2], 3, (255, 0, 0), 1)
199       cv2.circle(combinedoutliers, pt2[0:2], 3, (0, 255, 0), 1)
200       cv2.line(combinedoutliers, pt1[0:2], pt2[0:2], (0, 0, 255), 1)
201
202
203 def generatepanorama(H, images, path):
204   images = [img.copy() for img in images]
205   N = len(images)
206   mid = N // 2
207
208   midH = np.eye(3)
209   for idx in range(mid, len(H)):
210     midH = np.matmul(midH, np.linalg.inv(H[idx])) #H23, H34
211     H[idx] = midH
212
213   midH = np.eye(3)
214   for idx in range(mid - 1, -1, -1):
215     midH = np.matmul(midH, H[idx]) #H12, H01
216     H[idx] = midH
217
218   H.insert(mid, np.eye(3))
219   tx = 0
220   for idx in range(mid):
221     tx += images[idx].shape[1]
222   translation = np.array([[1, 0, tx], [0, 1, 0], [0, 0, 1]], dtype=float)
223
224   height, width = 0, 0
225   for idx in range(N):
226     height = max(height, images[idx].shape[0])

```

```

227     width += images[idx].shape[1]
228
229     combined = np.zeros((height, width, 3), np.uint8)
230     for idx in range(N):
231         Hcurr = np.matmul(translation, H[idx])
232         combined = pixelmapping(combined, images[idx], Hcurr)
233
234     cv2.imwrite(path + "panorama.jpg", combined)
235
236 def pixelmapping(combined, image, H):
237     height, width,   = image.shape
238     combinedheight, combinedwidth,   = combined.shape
239     H = np.linalg.inv(H)
240
241     Xarray, Yarray = np.meshgrid(np.arange(0, combinedwidth, 1), np.arange(0,
242                                   combinedheight, 1))
242     pixels = np.vstack((Xarray.ravel(), Yarray.ravel())).T
243     pixels = np.hstack((pixels[:, 0:2], pixels[:, 0:1]*0+1))
244
245     transformedpixel = np.dot(H, pixels.T)
246     transformedpixel = transformedpixel / transformedpixel[2, :]
247     transformedpixel = transformedpixel.T[:, 0:2].astype('int')
248
249     validpixels, validtransformed = findvalidpixels(pixels,
250     transformedpixel, width - 1, height - 1)
251     for idx in range(validpixels.shape[0]):
252         if((combined[validpixels[idx, 1], validpixels[idx, 0]] != 0).all()
253 == False):
254             combined[validpixels[idx, 1], validpixels[idx, 0]] = image[
255             validtransformed[idx, 1], validtransformed[idx, 0]]
256
257     return combined
258
259 def findvalidpixels(pixels, transformed, width, height):
260     xMin = transformed[:,0] >= 0
261     transformed = transformed[xMin, :]
262     pixels = pixels[xMin, :]
263
264     xMax = transformed[:,0] <= width
265     transformed = transformed[xMax, :]
266     pixels = pixels[xMax, :]
267
268     yMin = transformed[:,1] >= 0
269     transformed = transformed[yMin, :]
270     pixels = pixels[yMin, :]
271
272     yMax = transformed[:,1] <= height
273     transformed = transformed[yMax, :]
274     pixels = pixels[yMax, :]
275
276 if name == "main":

```

```

277 # Task 1
278 images1 = getimages(r"hw05/Images", task=1)
279 optimumH, bestinliers = ransac(images1, task=1, LM=False)
280 levenmarq = "withoutLM"
281 generatepanorama(optimumH, images1, path="hw05/panaroma/Task"+str(1)+"/"
+ levenmarq)

282
283 optimumH, bestinliers = ransac(images1, task=1, LM=True)
284 levenmarq = "LM"
285 generatepanorama(optimumH, images1, path="hw05/panaroma/Task"+str(1)+"/"
+ levenmarq)

286 # Task 2
287 images2 = getimages(r"hw05/Images", task=2)
288 optimumH, bestinliers = ransac(images2, task=2, LM=False)
289 levenmarq = "withoutLM"
290 generatepanorama(optimumH, images2, path="hw05/panaroma/Task"+str(2)+"/"
+ levenmarq)

291
292 optimumH, bestinliers = ransac(images2, task=2, LM=True)
293 levenmarq = "LM"
294 generatepanorama(optimumH, images2, path="hw05/panaroma/Task"+str(2)+"/"
+ levenmarq)

```

Listing 1: The Source Code

2.3 Inputs and Outputs



Figure 1: First Set of Inputs

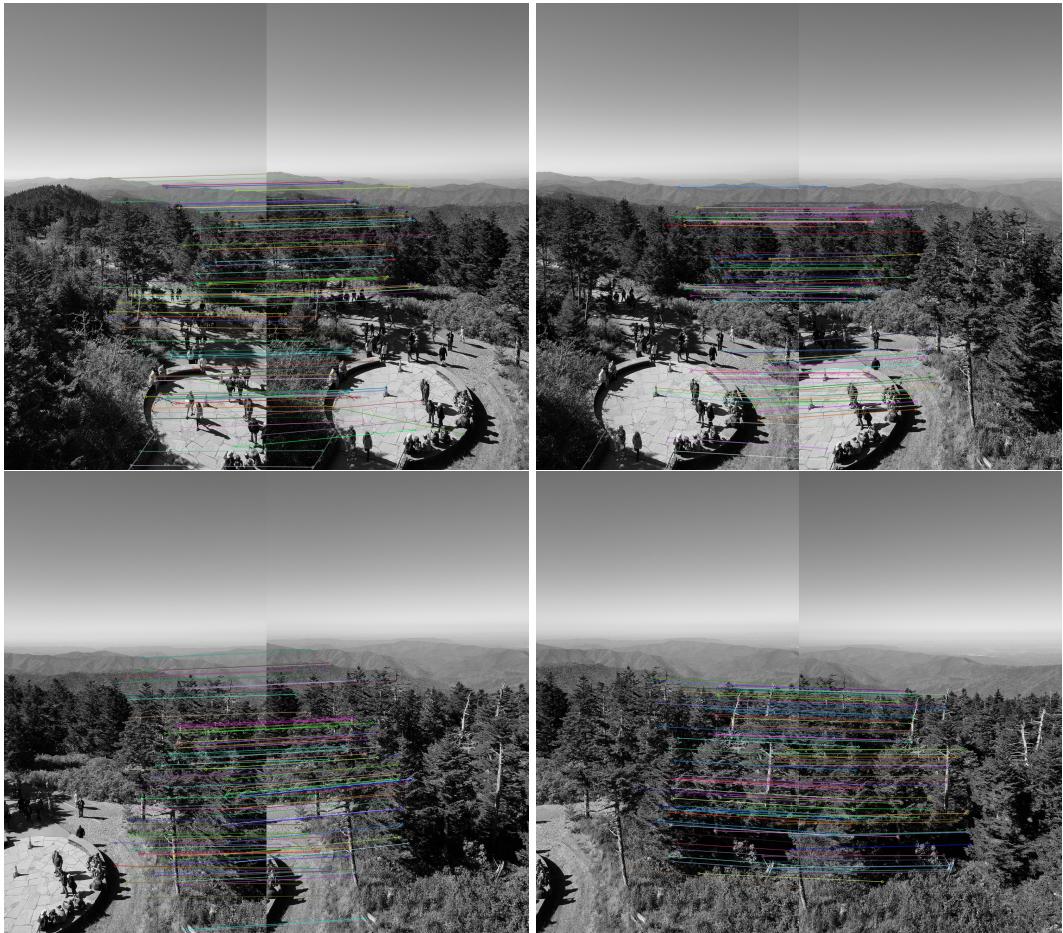


Figure 2: SIFT Correspondences for Input Set 1

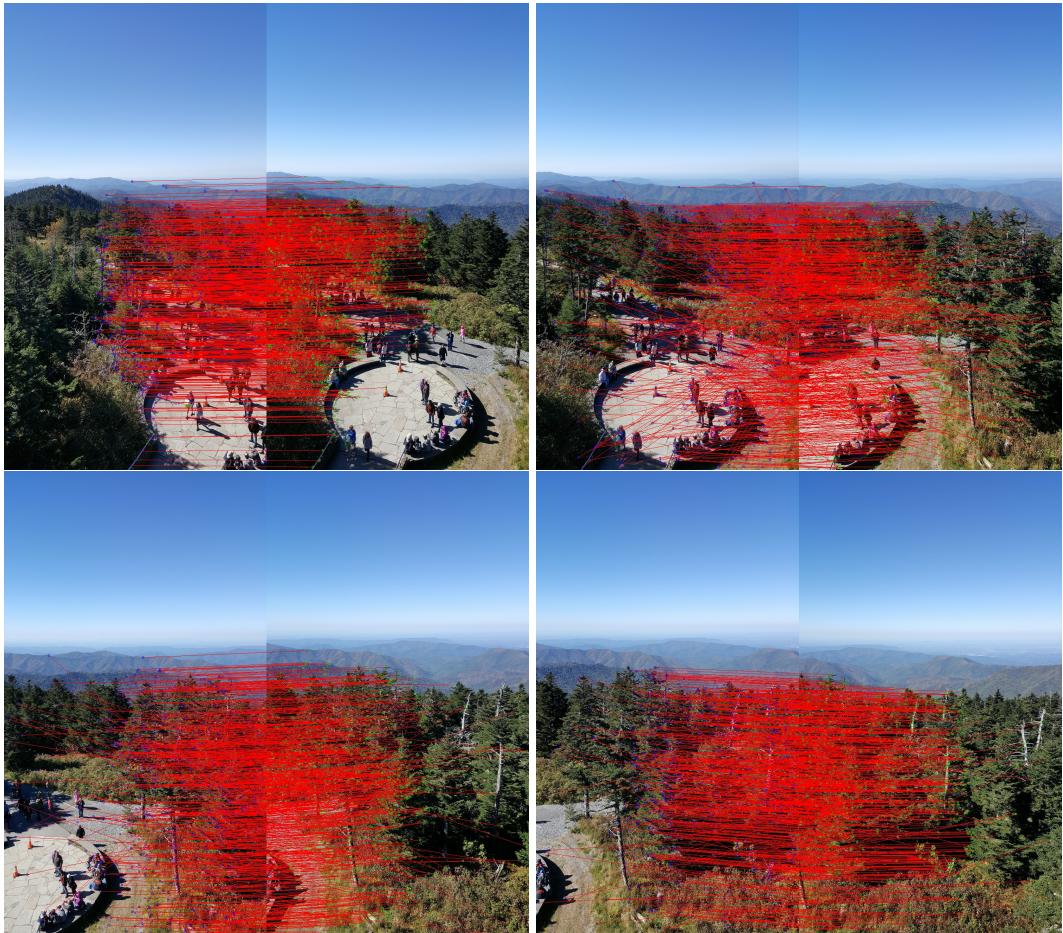


Figure 3: Inliers for Input Set 1

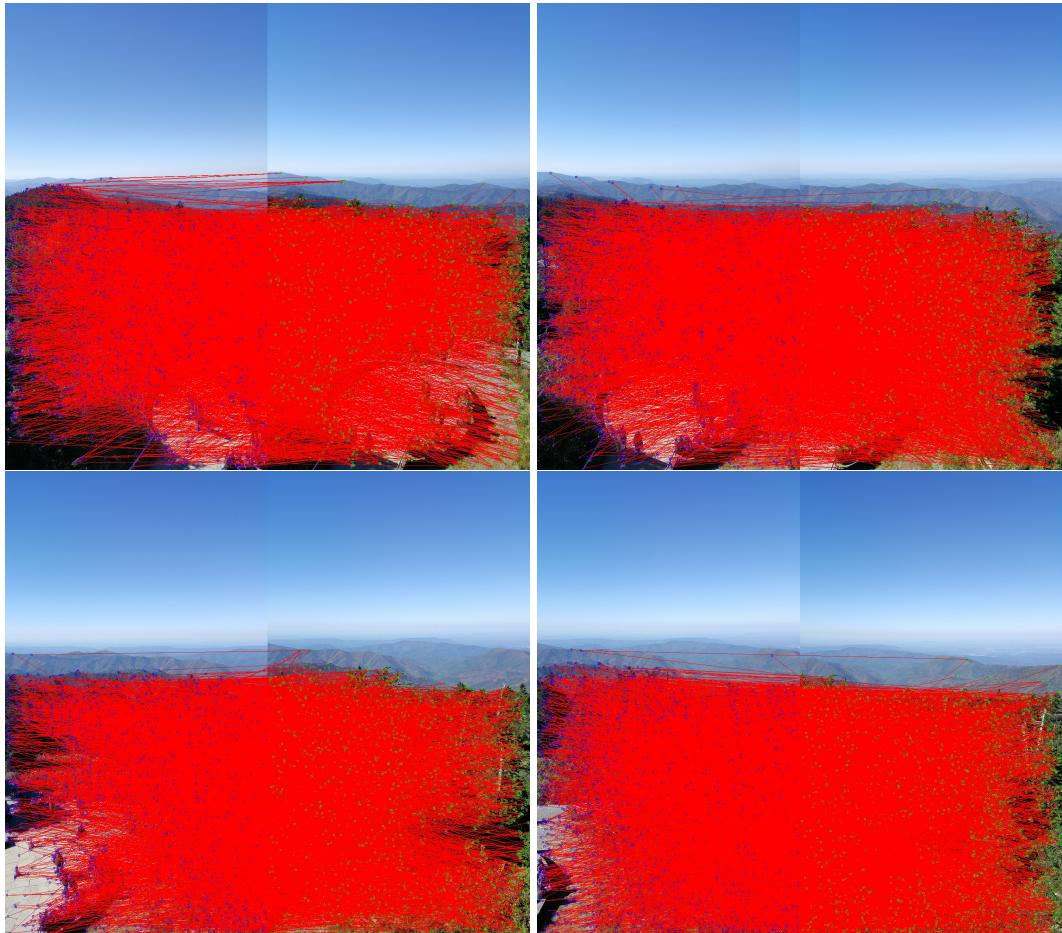


Figure 4: Outliers for Input Set 1



(a) With Levenberg-Marquadt



(b) Without Levenberg-Marquadt

Figure 5: Panoramas for Input Set 1



Figure 6: Second Set of Inputs

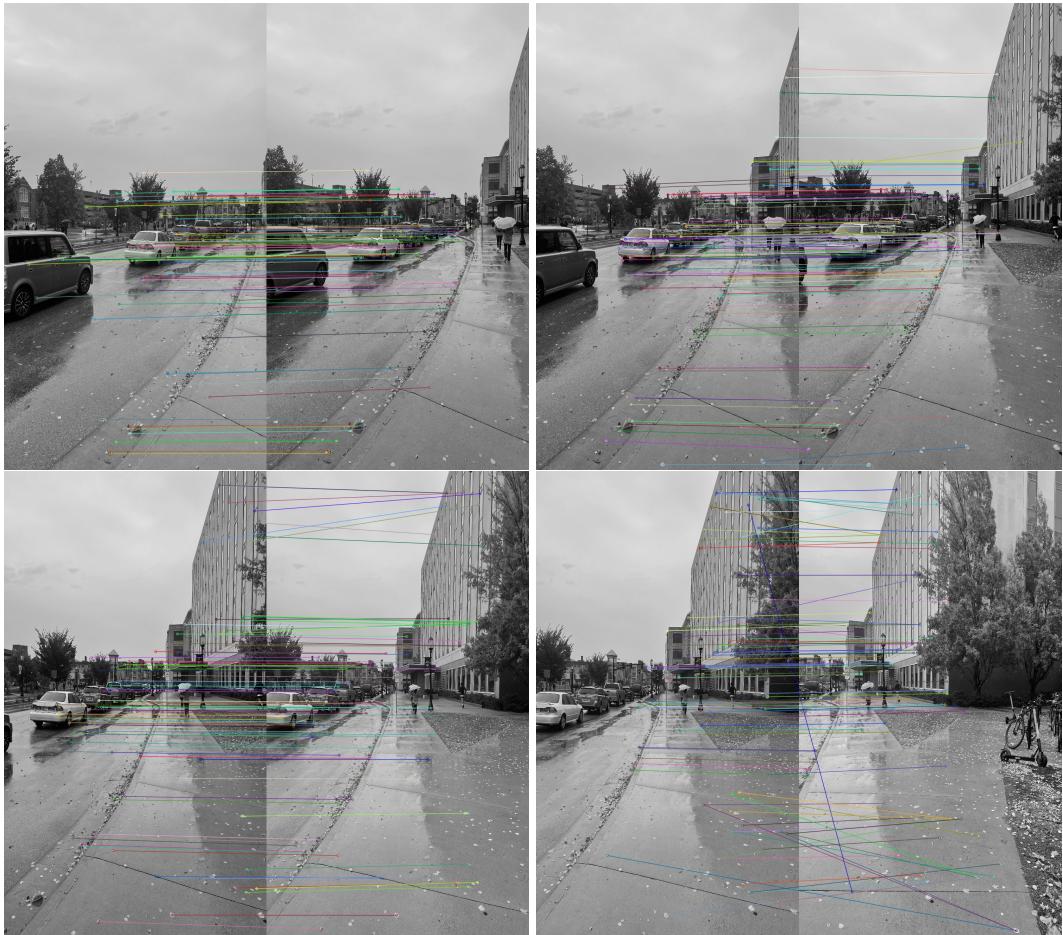


Figure 7: SIFT Correspondences for Input Set 2

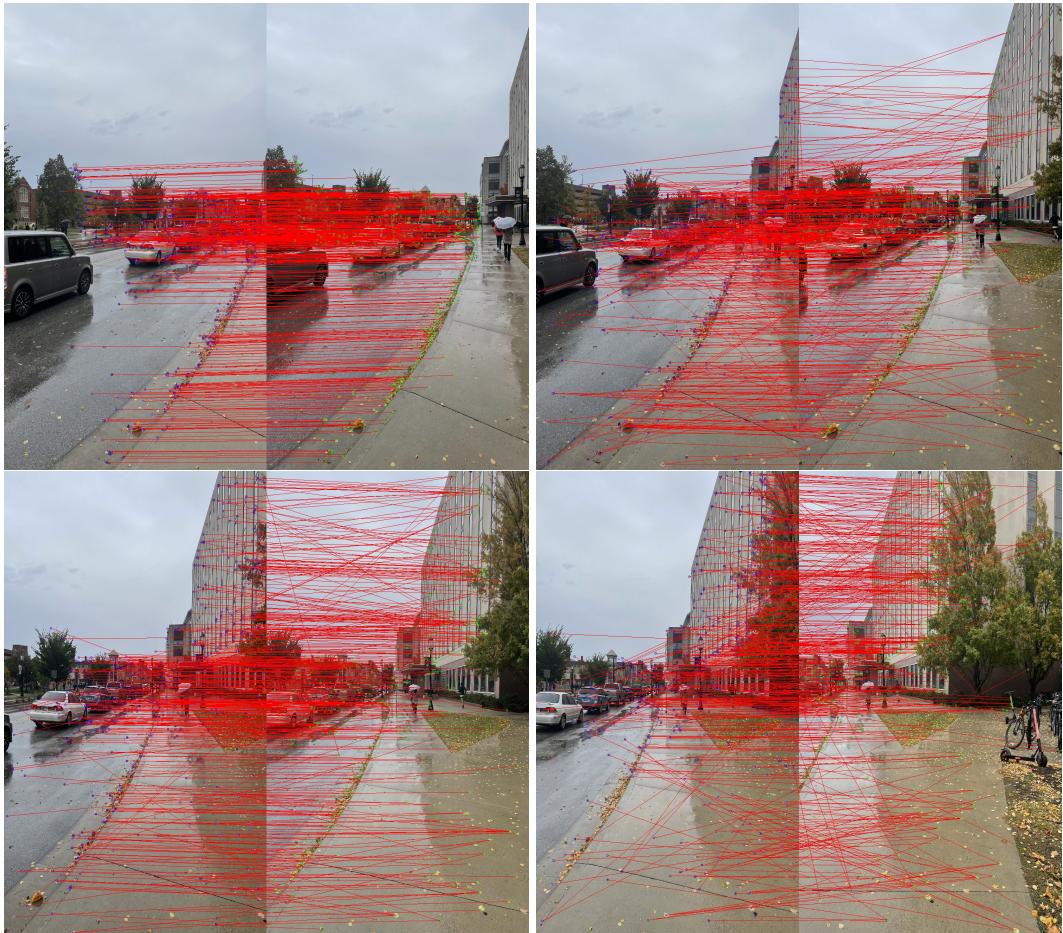


Figure 8: Inliers for Input Set 2

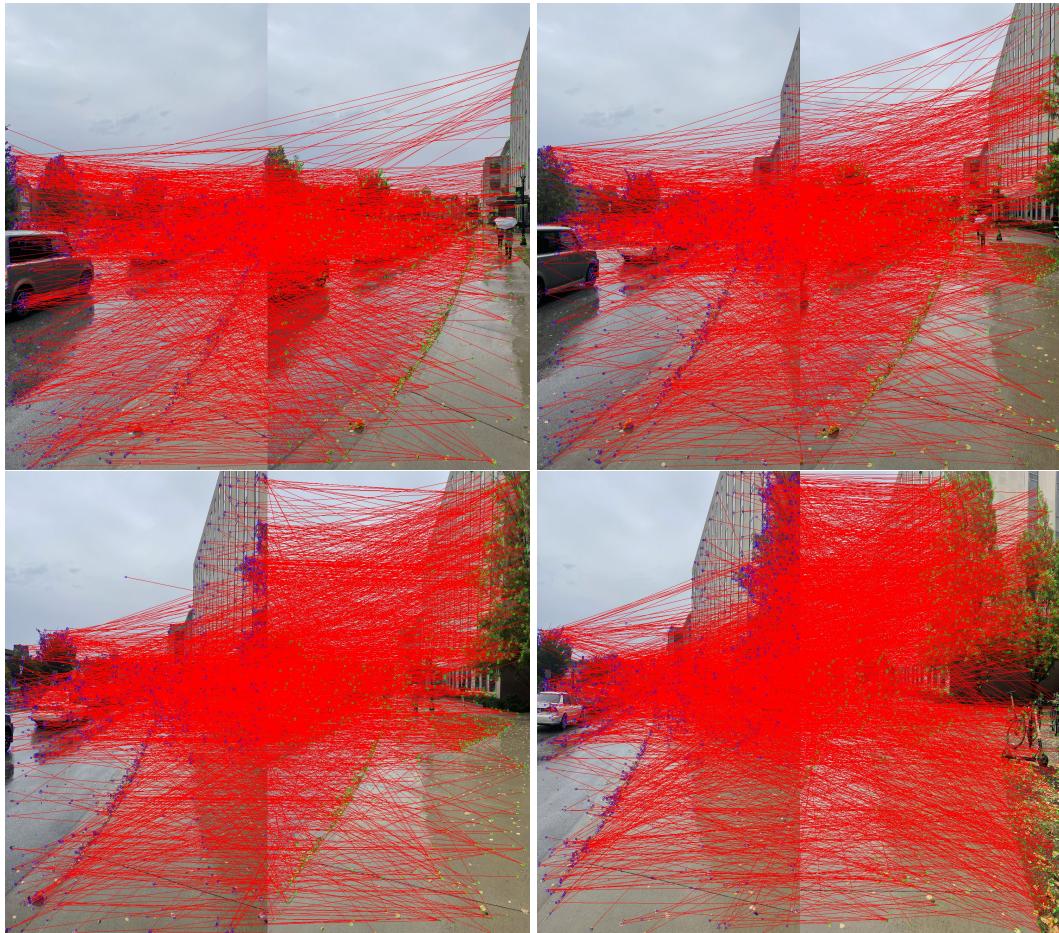


Figure 9: Outliers for Input Set 2



(a) With Levenberg-Marquadt



(b) Without Levenberg-Marquadt

Figure 10: Panoramas for Input Set 2