

1 Theory

1.1 Projective Stereo Reconstruction

The goal of projective stereo reconstruction is to construct a 3D model of a scene from its two views. There are many steps involved in making this possible which will be explained below.

1.1.1 Image Rectification

Image rectification involves transforming a pair of stereo images (so a left image and right image of the same scene) such that the pixels corresponding to the world coordinate in the left image is mapped to the same row as the corresponding world coordinates in the right image.

The process of image rectification involves”

- 1) We first normalize the eight coordinates from both stereo images. This is because the 8-point algorithm of constructing the fundamental matrix requires a translation and scaling of each image so that its centroid is at the origin and the RMS distance of the points from the origin is equal to $\sqrt{2}$ [1]. This is accomplished by first

- I. Compute the mean of the x-coordinates and y-coordinates
- II. Calculate the distance of each x-coordinate and y-coordinate to its respective mean coordinate
- III. Compute the mean of the distance vector calculated in step (II), denoted \bar{D}
- IV. Let $c = \frac{\sqrt{2}}{\bar{D}}$
- V. Create the normalization matrix such that the RMS is $\sqrt{2}$ as shown below

$$T = \begin{bmatrix} c & 0 & -cx \\ 0 & c & -cy \\ 0 & 0 & 1 \end{bmatrix} \quad (1)$$

- VI. The normalized pixel is therefore a matrix multiplication of the homogeneous coordinates and T
- 2) Calculate the initial estimate of the fundamental matrix F . This is accomplished by the following process
 - I. Using the normalized coordinates of the left image (\hat{x}, \hat{y}) and the normalized coordinates of the right image (\hat{x}', \hat{y}') to construct the 8×9 matrix A where each row is as shown below

$$A[\text{idx}] = [\hat{x}\hat{x}' \ \hat{y}\hat{x}' \ \hat{x}' \ \hat{x}\hat{y}' \ \hat{y}'\hat{y} \ \hat{y}' \ \hat{x} \ \hat{y} \ 1] \quad (2)$$

- II. We then solve the equation $Af = 0$ where f is a one-dimensional vector containing all the eight unknowns of F . This is done by doing a singular value decomposition (SVD) of A to get the initial estimate of F
- III. The fundamental matrix must be conditioned to enforce the requirement: $\text{rank}(F) = 2$ to construct a linear least squares solution. This is done by taking the SVD of F , and then setting the smallest eigenvalue of D to zero, we denote this as D' . Once we get D' , we do a matrix multiplication of the u from the SVD of F , with D' , and with vh from the SVD of F : $F_{\text{conditioned}} = u \times D' \times vh$
- IV. Finally, we de-normalize the conditioned initial estimate of F by doing a matrix multiplication of the normalization matrix T of both stereo images on F . In other words: $T_R^T \times F \times T_L$
- 3) We then optimize the denormalized F by using a non-linear least-squares optimization method called Levenberg-Marquadt algorithm to minimize the error to ensure proper rectification of the images. The cost function is as follows

$$\text{cost} = \sum_i (\|x_i - \hat{x}_i\|^2 + \|x'_i - \hat{x}'_i\|^2) \quad (3)$$

- 4) Using the optimized fundamental matrix, we then compute the epipole lines by solving these two following equations

$$F\vec{e} = 0 \quad (4) \qquad e^{\vec{T}}F = 0 \quad (5)$$

- 5) Using the canonical camera configuration method, we compute the projection matrices P and P' as follows

$$P = [I|0] = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \quad (6)$$

$$P' = [[\vec{e}']_x F | \vec{e}'] \text{where } [\vec{e}']_x = \begin{bmatrix} 0 & -e'_z & e'_y \\ e'_z & 0 & -e'_x \\ -e'_y & e'_x & 0 \end{bmatrix} \quad (7)$$

- 6) Finally, we need to determine a homography $H_{\{L,R\}}$ that would rectify the stereo images. First, we need to calculate H_R using the following method

I. Create a translation matrix to translate the image such that the center is at the origin

$$T = \begin{bmatrix} 1 & 0 & \frac{-w}{2} \\ 0 & 1 & \frac{-h}{2} \\ 0 & 0 & 1 \end{bmatrix} \quad (8)$$

II. Create the rotation matrix with an angle $\theta = -\arctan \frac{e'_y - h/2}{e'_x - w/2}$

$$R = \begin{bmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \quad (9)$$

III. Create a matrix G that would send the epipoles to infinity with a scale factor $f = (e'_x - \frac{w}{2}) \cos \theta - (e'_y - \frac{h}{2}) \sin \theta$

$$\begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ -\frac{1}{f} & 0 & 1 \end{bmatrix} \quad (10)$$

IV . Rectify the right image center by first calculating the homography $H_{R\text{-center}} = GRT$ and then multiplying it with the centroid of the image. The rectified center has coordinates (\tilde{x}, \tilde{y})

V. Create a second translation matrix that would move the rectified image center back to the original image center

$$T_2 = \begin{bmatrix} 1 & 0 & w/2 - \tilde{x} \\ 0 & 1 & h/2 - \tilde{y} \\ 0 & 0 & 1 \end{bmatrix} \quad (11)$$

VI. The final homography to rectify the right image is given by $H_R = T_2 GRT$

- 7) To calculate H_L we use the following method

- I. Repeat steps I-VI from the method used to compute the rectification homography for the right image
- II. H_L is obtained by the matrix that minimizes the least square distance between the transformed points [1]

$$\text{cost} = \sum_i (ax_i + by_i + c - x_i^2) \quad (12)$$

III. The constants from the minimized cost function will be used to construct the homography H_a

$$H_a = \begin{bmatrix} a & b & c \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \quad (13)$$

IV. Using H_a we repeat steps V-VI from the method to compute H_R

- 8) Using H_L and H_R rectify the left and right stereo images

1.1.2 Interest Point Detection

To make a 3D reconstruction of the stereo images, we need many correspondences between the two rectified images. We get these correspondences by

- 1) Perform edge detection on the rectified images to obtain key points
- 2) The corresponding points exist on the same row of the rectified images
- 3) SSD or NCC is used to filter our correspondences

1.1.3 3D Projective Reconstruction

To compute the scene structure from a set of correspondences for a given camera pair (P_L, P_R) and stereo correspondences $\vec{X} = (x, y)$ and $\vec{X}' = (x', y')$ by constructing matrix A

$$A = \begin{bmatrix} x\vec{P}_3^T - \vec{P}_1^T \\ y\vec{P}_3^T - \vec{P}_2^T \\ x'\vec{P}'_3^T - \vec{P}'_1^T \\ y'\vec{P}'_3^T - \vec{P}'_2^T \end{bmatrix} \quad (14)$$

We then solve the equation $A\vec{X} = 0$ using singular value decomposition.

1.1.4 Inputs



Figure 1: Stereo Images

1.1.5 Outputs



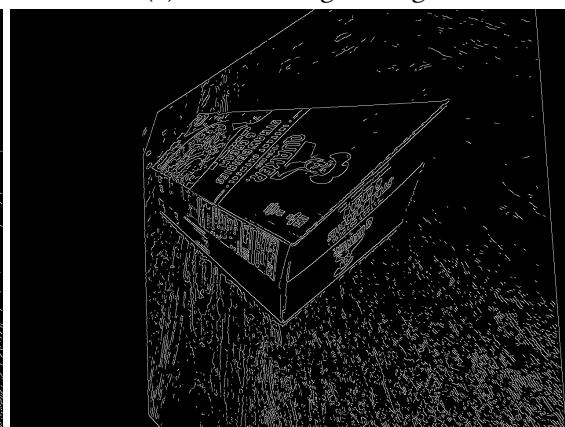
(a) Rectified Left Image



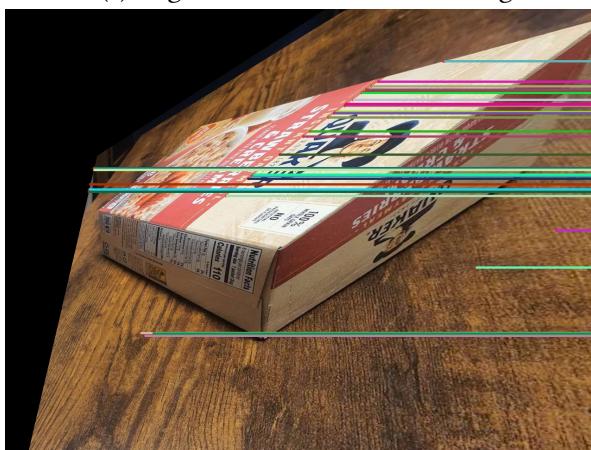
(b) Rectified Right Image



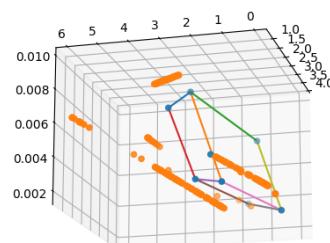
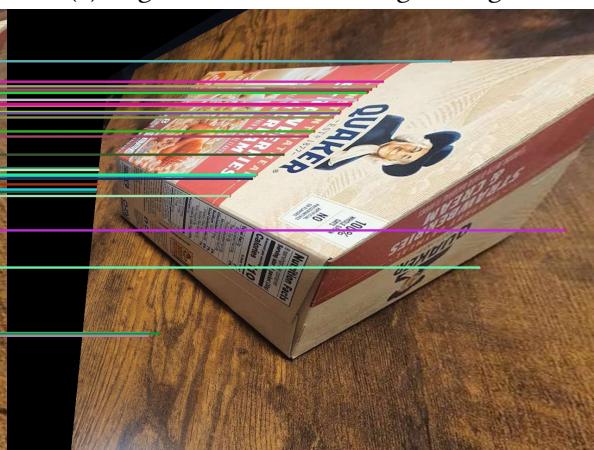
(c) Edges of the Rectified Left Image



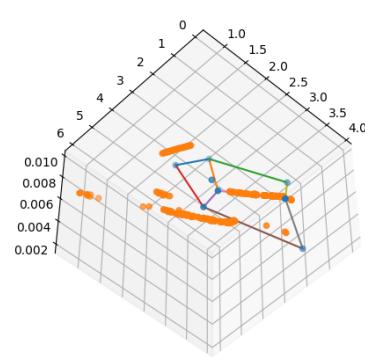
(d) Edges of the Rectified Right Image



(e) Correspondences



(f) 3D Reconstruction View 1



(g) 3D Reconstruction View 2

1.2 Loop and Zhang Algorithm

The Loop and Zhang algorithm decomposes the rectifying homographies:

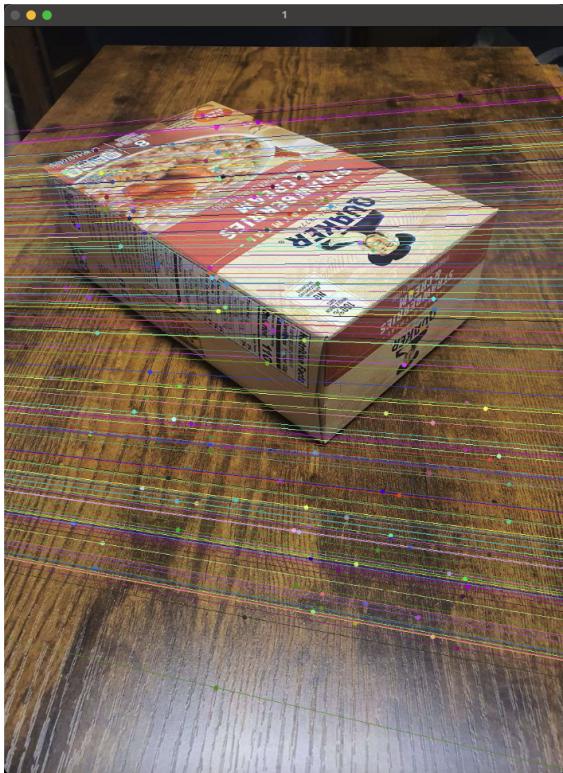
$$H_L = H_{\text{sh}} H_{\text{sim}} H_p \quad (15)$$

$$H_R = H'_{\text{sh}} H'_{\text{sim}} H'_p \quad (16)$$

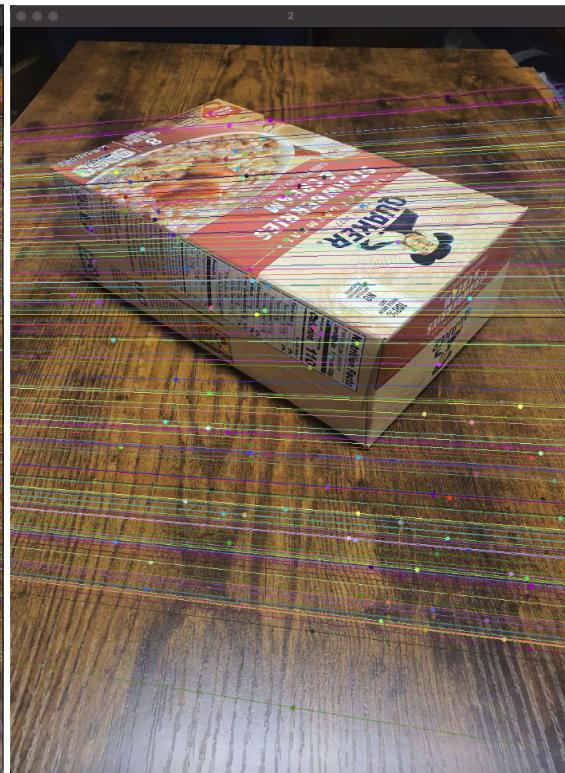
where H_p and H'_p are homographies that will send the epipoles to infinity in their respective image planes. The direction in which the epipoles will go to infinity is determined by the direction that causes the minimal amount of distortion. H_{sim} and H'_{sim} are similarity homographies which can rotate, translate, and uniformly scale an image. The purpose of the similarity homography is to rotate the epipoles onto the x-axis without causing any distortion. Finally, H_{sh} and H'_{sh} are shearing homographies. Since a non-linear distortion of the two images caused by projective homography cannot be undone by similarity transformations, the best we can do is reduce the distortion by introducing additional degrees of freedom into the overall rectification transformation which is where the shearing homographies come into the equation [2].

Despite the homographies being homogeneous, setting the last component of the matrices to one will prevent sending the origin in the image I to infinity. So typically, the corner of the sensor array is used as the image origin and the likelihood of the epipole coinciding with such an origin will be slim. The two cameras need to be angled sharply toward each other for this to happen [2].

1.2.1 Outputs



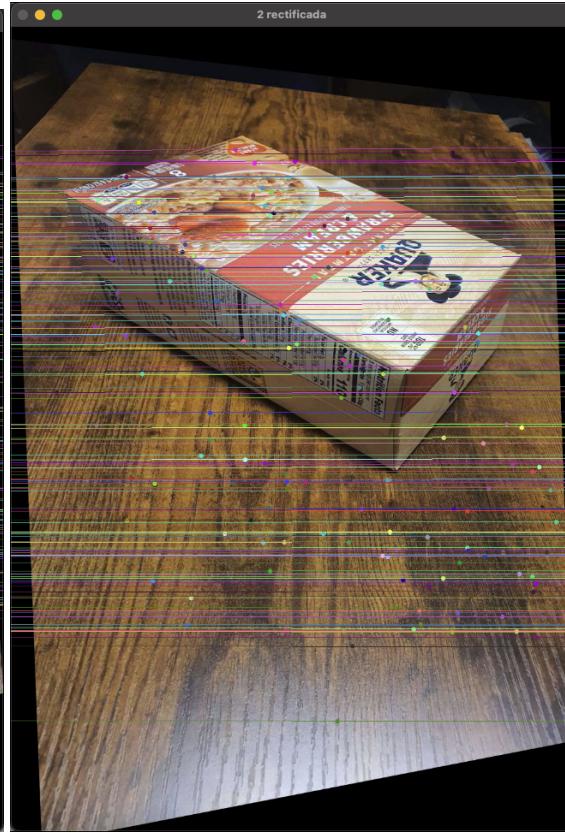
(a) Left Image



(b) Right Image



(c) Rectified Left Image



(d) Rectified Right Image

Figure 3: Loop and Zhang Outputs

1.2.2 Observations

The rectified images obtained from the Loop and Zhang algorithm are different from the pipeline I developed in task 1. Just like the correspondence lines from task 1, the correspondence lines from the Loop and Zhang algorithms are completely horizontal suggesting that the images have been successfully rectified.

1.3 Dense Stereo Matching

Dense Stereo Matching is similar to the canny edge detector method to correspondences between stereo images. This method is better than the canny edge detector because it does not return a subset of correspondences but rather it can match all the pixels. The Dense Stereo Matching is conducted by doing a census transform which is done as follows

- 1) Calculate the d_{\max} from the ground truth disparity map. The d_{\max} obtained for this homework is 52.
- 2) For a pixel in the left image, save the intensity of the pixels around it in a $M \times M$ neighborhood.
- 3) Create a $M \times M$ window for the right image around the pixel with the same coordinates as the one in the left image and save the intensity of the pixels in this window for d_{\max} times
- 4) Pixels with intensity greater than the intensity of the center pixel are assigned 1, otherwise it is assigned a 0
- 5) Bitwise XOR the two binary neighborhoods obtained from step (3)
- 6) The cost associated to a pixel in the right image is the number of 1s obtained from step (4)
- 7) The pixel at a specific row and column in the left image will be assigned the lowest cost calculated in the previous step on the disparity map

1.3.1 Inputs

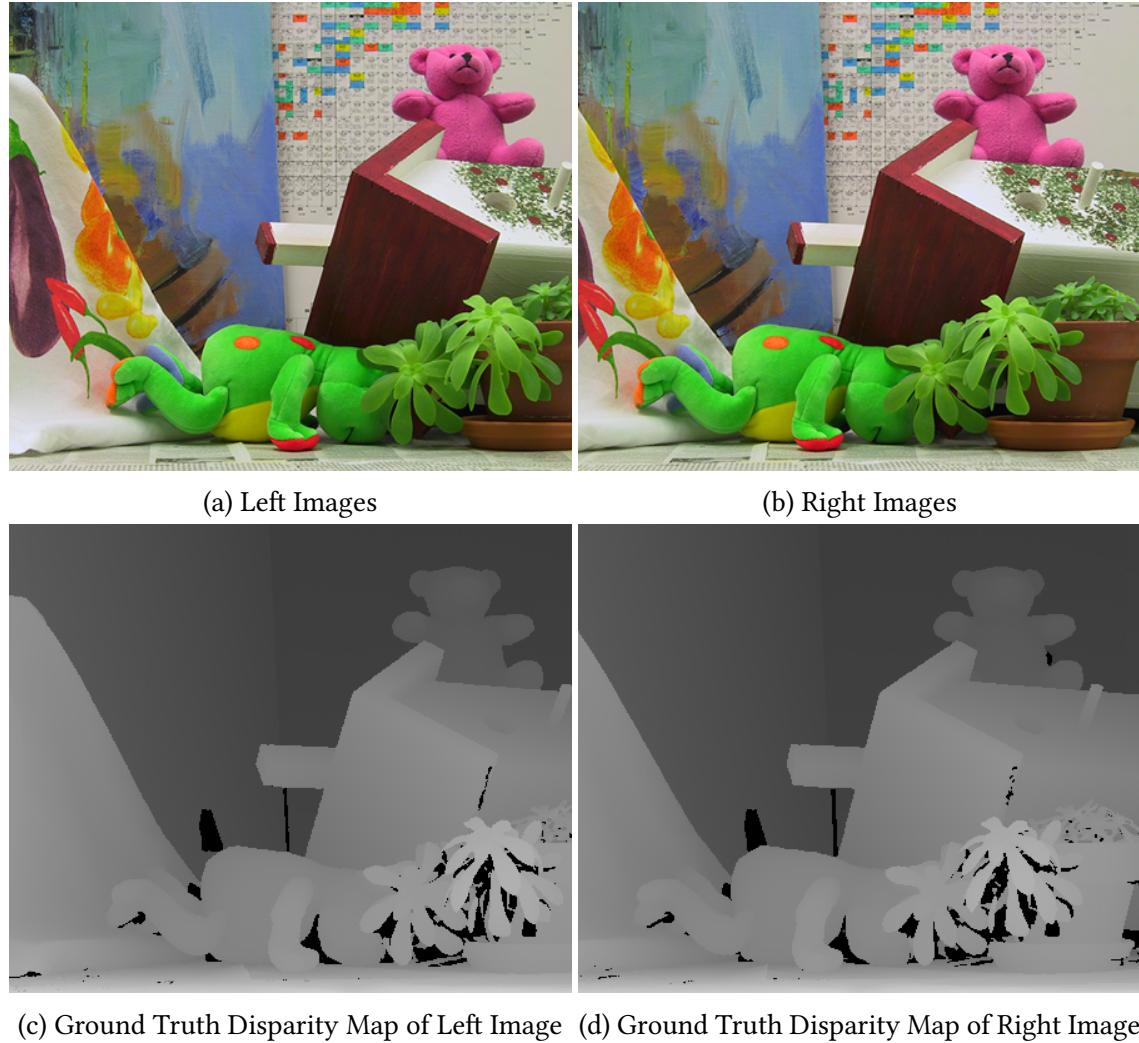


Figure 4: Dense Stereo Matching Inputs

1.3.2 Outputs

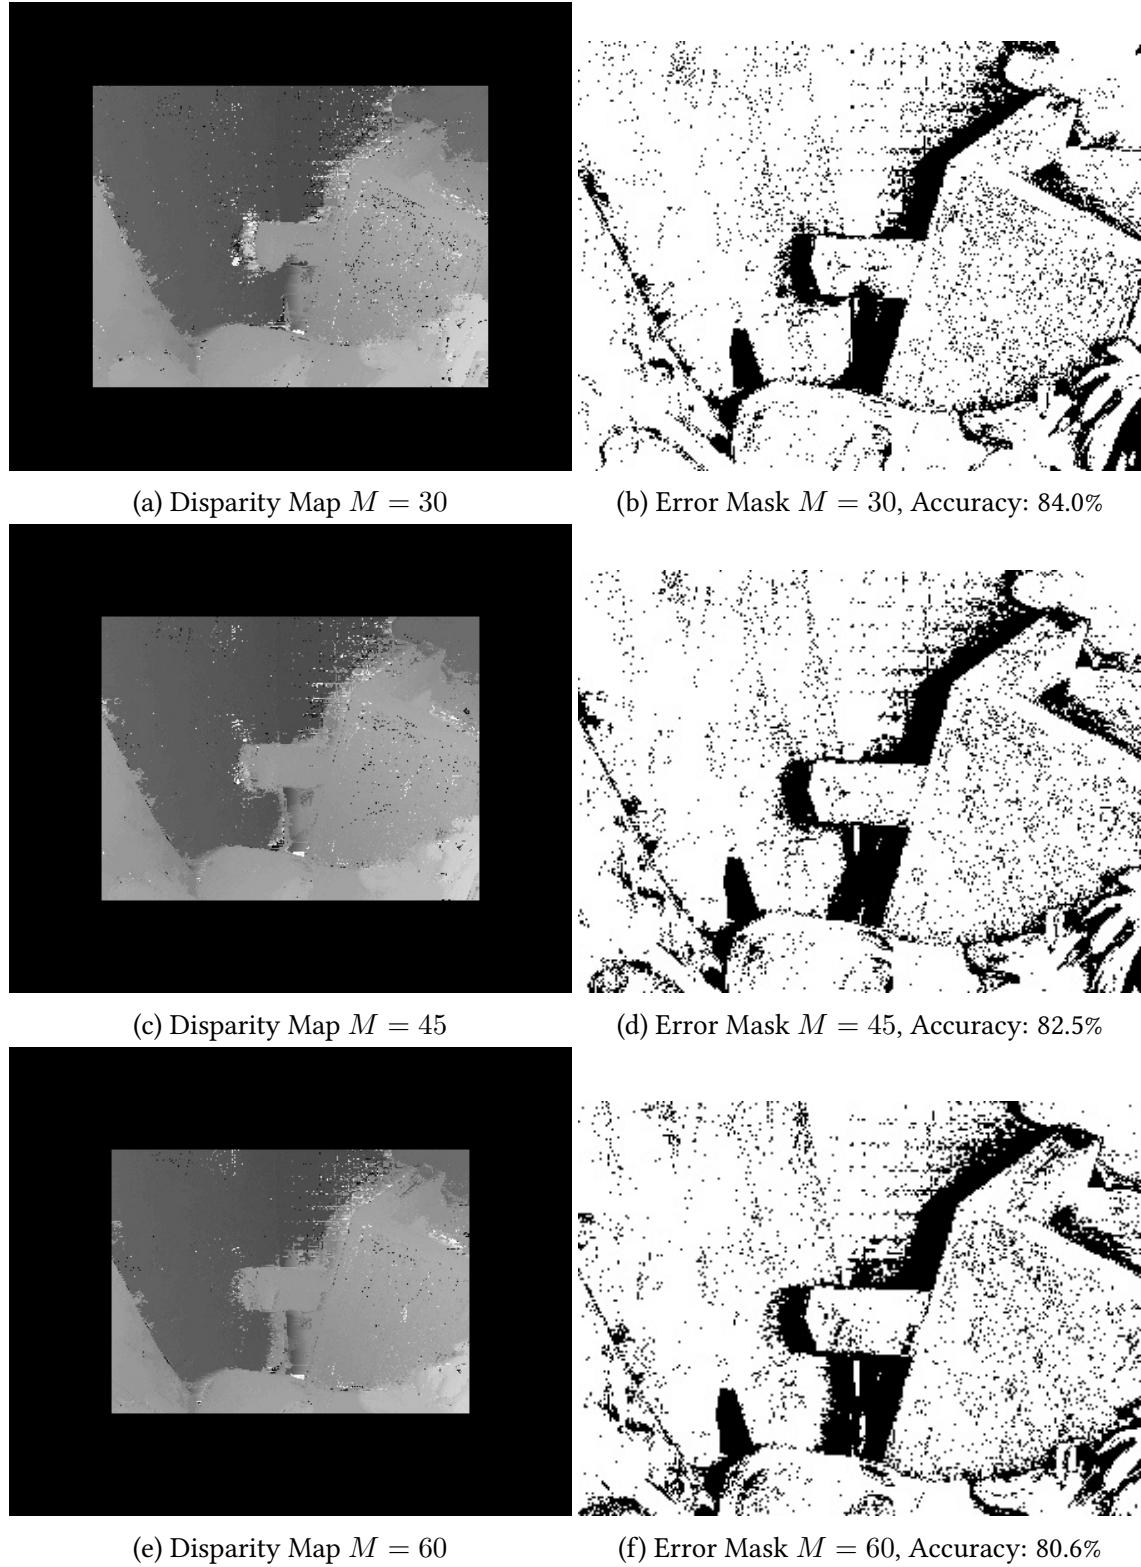


Figure 5: Dense Stereo Matching Outputs

1.3.3 Observations

As the size of the window M increased, the disparity map became a lot more clearer but the quality of the error masks and accuracy kept decreasing.

2 Source Code

```
1 # Name: Nikita Ravi
2 # Class: ECE 66100
3 # Homework #9
4 # Deadline: 11/27/2022
5
6 # Import Modules
7 import cv2
8 import numpy as np
9 import math
10 import random
11 import matplotlib.pyplot as plt
12 from scipy.optimize import leastsquares
13
14 def displayimage(image, points=True):
15     def clickevent(event, x, y, flags, params):
16         # This function was inspired by https://www.geeksforgeeks.org/displaying-the-coordinates-of-the-points-clicked-on-the-image-using-python-opencv/
17         if(event == cv2.EVENTLBUTTONDOWN):
18             print(x, y)
19             font = cv2.FONT_HERSHEY_SIMPLEX
20             cv2.circle(image, (x, y), 1, (0,0,255), thickness=-1)
21             cv2.putText(image, str(x) + ', ' +
22                         str(y), (x,y), font,
23                         1, (0, 0, 255), 2)
24             cv2.imshow('window', image)
25
26
27         cv2.imshow("window", image)
28     if(points):
29         cv2.setMouseCallback('window', clickevent)
30         cv2.waitKey(0)
31         cv2.destroyAllWindows()
32         quit()
33
34     def drawcircles(path, image, pts):
35         for idx, pt in enumerate(pts):
36             x, y = int(pt[0]), int(pt[1])
37             cv2.circle(image, (x, y), radius=6, color=(0,0,255), thickness=-1)
38
39         cv2.imwrite(path, image)
40
41     def normalizepoints(correspondences):
42         """
43             Normalization is necessary for the 8-point algorithm of constructing the
44             fundamental matrix to be successful. The suggested normalization
```

```

44     is a translation and scaling of each image so that the centroid of the
45     reference points is at the origin of the coordinates and the RMS
46     distance of the points from the origin is equal to sqrt(2) (Hartley and
47     Zisserman, 2003).
48     """
49
50     xcoord = correspondences[:, 0]
51     ycoord = correspondences[:, 1]
52
53     # Compute the mean distance
54     xmean, ymean = np.mean(xcoord), np.mean(ycoord)
55     distances = np.sqrt((xcoord - xmean) ** 2 + (ycoord - ymean) ** 2)
56     meandist = np.mean(distances)
57
58     # Create normalizing transformation matrix consisting of translation and
59     # scaling
60     c = np.sqrt(2) / meandist
61     tMat = np.array([[c, 0, -c*xmean],
62                      [0, c, -c*ymean],
63                      [0, 0, 1]])
64
65     # Apply the translation matrices
66     normalizedpts = np.hstack((correspondences, np.ones((len(correspondences),
67                                         1)))) # Make it homogeneous
68     normalizedpts = np.dot(tMat, normalizedpts.T).T
69
70     return normalizedpts, tMat
71
72
73 def conditionF(F):
74     # Fundamental matrix F must be conditioned to enforce the requirement rank(F)
75     # = 2 to construct a linear least squares solution
76     u, d, vh = np.linalg.svd(F)
77     dprime = np.array([[d[0], 0, 0],
78                        [0, d[1], 0],
79                        [0, 0, 0]])
80     conditionedF = u @ dprime @ vh
81     return conditionedF
82
83
84 def calculatefundamentalmatrix(normalizedptsL, normalizedptsR):
85     A = np.zeros((8, 9))
86     for corr in range(len(normalizedptsL)):
87         x1, y1 = normalizedptsL[corr, 0], normalizedptsL[corr, 1]
88         x2, y2 = normalizedptsR[corr, 0], normalizedptsR[corr, 1]
89
90         A[corr] = np.array([x1 * x2, x2 * y1, x2, y2 * x1, y2 * y1, y2,
91                           x1, y1, 1])
92
93     u, s, vh = np.linalg.svd(A)
94     Fflatten = vh[-1]
95     F = np.reshape(Fflatten, (3, 3))
96     return conditionF(F)
97
98
99 def computeEx(e):
100    Ex = np.array([[0, -e[2], e[1]],
101                  [e[2], 0, -e[0]]],
102
```

```

92         [-e[1], e[0], 0]])
93
94     return E x
95
96 def findepipoles(F):
97     u,s,vh = np.linalg.svd(F)
98     eL = vh[-1, :].T
99     eL = eL / eL[2] # Make it homogeneous
100
101    eR = u[:, -1]
102    eR = eR / eR[2] # Make it homogeneous
103
104    E x = computeEx(eR)
105    return eL, eR, E x
106
107 def computeprojectionmatrices(F, eR, s):
108     P1 = np.hstack((np.eye(3), np.zeros((3,1))))
109     P2 = np.hstack((np.dot(s, F), np.transpose([eR])))
110
111     return P1, P2
112
113 def costfunction(f, ptsL, ptsR):
114     F = np.reshape(f, (3, 3))
115     eL, eR, E x = findepipoles(F)
116     P1, P2 = computeprojectionmatrices(F, eR, E x)
117
118     ptsL = np.hstack((ptsL, np.ones((len(ptsL), 1)))) # Make it homogeneous
119     ptsR = np.hstack((ptsR, np.ones((len(ptsR), 1)))) # Make it homogeneous
120
121     dgeom = []
122     for corr in range(len(ptsL)):
123         A = np.zeros((4, 4)) # AX = 0
124         A[0] = ptsL[corr][0] * P1[2, :] - P1[0, :]
125         A[1] = ptsL[corr][1] * P1[2, :] - P1[1, :]
126         A[2] = ptsR[corr][0] * P2[2, :] - P2[0, :]
127         A[3] = ptsR[corr][1] * P2[2, :] - P2[1, :]
128
129         # Compute X
130         u, s, vh = np.linalg.svd(A)
131         X = vh[-1, :].T
132         Xhat = X / X[3] # Unit Vector of X
133         xL = np.dot(P1, Xhat)
134         xL = xL / xL[2] # Make it Homogeneous
135
136         xR = np.dot(P2, Xhat)
137         xR = xR / xR[2] # Make it Homogeneous
138
139         # Append the re-projection errors
140         dgeom.append(np.linalg.norm(xL - ptsL[corr]) ** 2)
141         dgeom.append(np.linalg.norm(xR - ptsR[corr]) ** 2)
142
143     return np.ravel(dgeom)
144
145 def computehomographyrectification(image, corrsL, corrsR, eL, eR):

```

```

146 height, width = image.shape[:2]
147
148 # Translation Matrix to translate the image such that the center is at the
149 # origin
150 T = np.array([[1, 0, -width / 2],
151               [0, 1, -height / 2],
152               [0, 0, 1]]) # Obtained Translation Matrix from https://web.stanford.
153 #edu/class/cs231a/coursenotes/03-epipolar-geometry.pdf
154
155 ##### Calculate HR or H' #####
156 # Rotation Matrix to move the epipoles towards the x-axis
157 theta = -np.arctan((eR[1] - height/2)/(eR[0] - width/2))
158 R = np.array([[math.cos(theta), -math.sin(theta), 0],
159               [math.sin(theta), math.cos(theta), 0],
160               [0, 0, 1]]) # Inspired by 2020 hw#1 solution
161
162 # Create matrix that would send the epipoles to infinity
163 f = (eR[0] - width / 2) * math.cos(theta) - (eR[1] - height / 2) * math.sin(
164     theta) # Inspired by 2020 hw#1 solution
165 G = np.array([[1, 0, 0],
166               [0, 1, 0],
167               [-1 / f, 0, 1]])
168
169 # Rectify the right image center
170 HRcenter = np.dot(G, np.dot(R, T))
171 Xprime = np.dot(HRcenter, np.array([width / 2, height / 2, 1])) # Apply
172     HRcenter to rectify the center of the right image (x', y')
173 T2 = np.array([[1, 0, width / 2 - Xprime[0]],
174               [0, 1, height / 2 - Xprime[1]],
175               [0, 0, 1]]) # Move the rectified center back to the true image
176     center
177
178 # Rectify the entire right image
179 HR = np.dot(np.dot(T2, G), np.dot(R, T))
180 HR = HR / HR[2, 2] # Make it Homogeneous
181
182 ##### Calculate HL or H #####
183 # Create a rotation matrix to move the epipole to the x-axis
184 theta = math.atan2(eL[1] - height / 2, (eL[0] - width / 2))
185 R = np.array([[math.cos(theta), -math.sin(theta), 0],
186               [math.sin(theta), math.cos(theta), 0],
187               [0, 0, 1]]) # Inspired by 2020 hw#1 solution
188
189 # Create matrix that would send the epipoles to infinity
190 f = math.cos(theta) * (eL[0] - width / 2) - math.sin(theta) * (eL[1] -
191     height / 2) # Inspired by 2020 hw#1 solution
192 G = np.array([[1, 0, 0],
193               [0, 1, 0],
194               [-1 / f, 0, 1]])
195
196 # Rectify the left image center
197 H0 = G @ R @ T
198
199 # HL is obtained by the matrix that minimizes the least squares fistance

```

```

194 # pg 307 (Hartley and Zisserman, 2003)
195 ptsL = np.hstack((corrsL, np.ones((len(corrsL), 1)))) # Make it Homogeneous
196 ptsR = np.hstack((corrsR, np.ones((len(corrsR), 1)))) # Make it Homogeneous
197
198 # Transform the points by their respective homographies
199 xL = np.dot(H0, ptsL.T).T
200 xR = np.dot(HR, ptsR.T).T
201
202 # Make them Homogeneous
203 # xL = xL / xL[:, 2]
204 xL[:, 0] = xL[:, 0] / xL[:, 2]
205 xL[:, 1] = xL[:, 1] / xL[:, 2]
206 xL[:, 2] = xL[:, 2] / xL[:, 2]
207
208 xR[:, 0] = xR[:, 0] / xR[:, 2] # Only need the x-coordinate as shown on pg.
209 # 307 to minimize the dist
210
211 # The least squares solution to an mxn system of equations Ax=b of rank n is
212 # given by x = (A+)b which minimizes ||x||. (pg. 590)
213 a,b,c = np.dot(np.linalg.pinv(xL), xR[:, 0])
214 Ha = np.array([[a, b, c], [0, 1, 0], [0, 0, 1]]) # Eq. 11.20 pg 306
215 HLcenter = np.dot(Ha, H0)
216 Xprime = np.dot(HLcenter, np.array([width / 2, height / 2, 1])) # Apply
217 # HRcenter to rectify the center of the right image (x', y')
218 Xprime = Xprime / Xprime[2] # Make it Homogeneous
219
220 T2 = np.array([[1, 0, width / 2 - Xprime[0]],
221                 [0, 1, height / 2 - Xprime[1]],
222                 [0, 0, 1]]) # Move the rectified center back to the true image
223 # center
224
225 # Rectify the entire image
226 HL = np.dot(T2, HLcenter)
227 HL = HL / HL[2, 2]
228
229 return HL, HR
230
231 def transform(domainimage, H):
232     height, width = domainimage.shape[:2]
233     return cv2.warpPerspective(domainimage, H, dsize=(height, width))
234
235 def detectedges(rectifiedimg, lowthreshold=60, highthreshold=70, edges=5):
236     gray = cv2.cvtColor(rectifiedimg, cv2.COLORBGR2GRAY) if len(rectifiedimg)
237     shape) == 3 else rectifiedimg
238     gray = cv2.GaussianBlur(gray, ksize=(7, 7), sigmaX=1.3)
239
240     return cv2.Canny(gray, threshold1=highthreshold, threshold2=lowthreshold,
241                      edges=edges)
242
243 def filtercorrespondences(imgL, imgR, corrs, kernelsize, maximumcorrs, mode
244 = "ncc"):
245     halfkernelsize = kernelsize // 2
246     if(maximumcorrs > len(corrs)):
247         print("Maximum correspondences exceeded the number of correspondences")

```

```

241 maximumcorrs = len(corrss)
242
243 dist = []
244 for corr in corrss:
245     corrL, corrR = corr
246     f1 = imgL[corrL[1] - halfkernelsize:corrL[1] + halfkernelsize + 1,
247                corrL[0] - halfkernelsize:corrL[0] + halfkernelsize + 1]
248
249     f2 = imgR[corrR[1] - halfkernelsize:corrR[1] + halfkernelsize + 1,
250                corrR[0] - halfkernelsize:corrR[0] + halfkernelsize + 1]
251
252 if(f1.shape != f2.shape):
253     continue
254
255 if(mode == "ssd"):
256     dist.append(np.sum((f1 - f2) ** 2))
257 elif(mode == "ncc"):
258     m1, m2 = np.mean(f1), np.mean(f2)
259     numerator = np.sum((f1-m1)*(f2-m2))
260     denominator = np.sqrt(np.sum((f1-m1)**2) * np.sum((f2-m2)**2))
261     distance = numerator / denominator
262     dist.append(1 - distance)
263
264 # Sort the correspondences and choose the best maxcorrespondence
265 sortedcorrespondences = [pt for d, pt in sorted(zip(dist, corrss), key=
266     lambda x: x[0])]
267 return sortedcorrespondences[0:maximumcorrs]
268
269 def getcolor():
270     r = random.randint(0, 255)
271     g = random.randint(0, 255)
272     b = random.randint(0, 255)
273     return (r,g,b)
274
275 def drawcorrespondences(imgL, imgR, width1, correspondences, path):
276     # From homework 4
277     combined = np.hstack((imgL, imgR))
278     for idx, corr in enumerate(correspondences):
279         if(not idx % 10):
280             pt1, pt2 = corr
281             cv2.circle(combined, (pt1[0], pt1[1]), radius=3, color=(0,0,255))
282             cv2.circle(combined, (pt2[0] + width1, pt2[1]), radius=3, color
283             =(0,0,255))
284             cv2.line(combined, pt1, (pt2[0] + width1, pt2[1]), color=getcolor(),
285             thickness=3)
286
287     cv2.imwrite(path, combined)
288
289 def detectcorrespondences(edgesL, edgesR, maxsearcharea):
290     grayEdgesL = cv2.cvtColor(edgesL, cv2.COLORBGR2GRAY) if len(edgesL.shape)
291     == 3 else edgesL
292     grayEdgesR = cv2.cvtColor(edgesR, cv2.COLORBGR2GRAY) if len(edgesR.shape)
293     == 3 else edgesR

```

```

290 listofcorrespondences = []
291 for row in range(grayEdgesL.shape[0]):
292     nonzeroidxL = np.where(grayEdgesL[row] != 0)[0] # Get indices of all the
293     edges that are not equal to 0
294     if(not np.size(nonzeroidxL)):
295         continue
296
297     # For each pixel in the left image, find the left-most pixel in the right
298     # image
299     for colL in nonzeroidxL:
300         searchkernel1 = edgesR[row, colL:colL + maxsearcharea + 1]
301         potentialnonzeroidxR = np.where(searchkernel1 != 0)[0]
302
303         if(not np.size(potentialnonzeroidxR)):
304             continue
305
306         colR = potentialnonzeroidxR[0] + colL
307         edgesR[row, colR] = 0 # Don't count this pixel anymore
308         listofcorrespondences.append([[colL, row], [colR, row]])
309
310 return listofcorrespondences
311
312 def projectivereconstruction(corr, PL, PR):
313     worldcoords = []
314
315     for idx, corr in enumerate(corr):
316         A = np.zeros((4, 4))
317         corrL, corrR = corr
318
319         A[0] = corrL[0] * PL[2, :] - PL[0, :]
320         A[1] = corrL[1] * PL[2, :] - PL[1, :]
321         A[2] = corrR[0] * PR[2, :] - PR[0, :]
322         A[3] = corrR[1] * PR[2, :] - PR[1, :]
323
324         u, s, vh = np.linalg.svd(A.T @ A)
325         vh = vh[-1, :].T
326         worldcoords.append(vh / vh[3])
327
328     worldcoords = np.reshape(worldcoords, (len(corr), 4))
329     return worldcoords
330
331 def transformpts(pts, H):
332     transformedpts = []
333     for pt in pts:
334         X = np.array([pt[0], pt[1], 1.])
335         Xprime = np.dot(H, X)
336         Xprime = Xprime / Xprime[-1]
337         transformedpts.append((Xprime[0], Xprime[1]))
338     return transformedpts
339
340 def projectivestereoreconstruction():
341     imgL = cv2.imread(r"/Users/nikitaravi/Documents/Academics/Year 4/Semester 2/
342     ECE 66100/hw09/Task3Images/imgL.JPG")

```

```

340 imgR = cv2.imread(r"/Users/nikitaravi/Documents/Academics/Year 4/Semester 2/
341 ECE 66100/hw09/Task3Images/imgR.JPG")
342
343 # Correspondence Points for the Left Image and Right Image
344 corrsL = np.int32(np.array([[66, 242],
345 [122, 417],
346 [549, 714],
347 [553, 532],
348 [846, 314],
349 [799, 496],
350 [354, 127],
351 [364, 424]]))
352
353 corrsR = np.int32(np.array([[71, 239],
354 [117, 420],
355 [611, 721],
356 [640, 531],
357 [883, 308],
358 [825, 492],
359 [355, 122],
360 [423, 430]]))
361
362 # Save the inputs with correspondence points
363 drawcircles(r"/Users/nikitaravi/Documents/Academics/Year 4/Semester 2/ECE
364 66100/hw09/stereoInputs/imgL.jpg", imgL.copy(), corrsL)
365 drawcircles(r"/Users/nikitaravi/Documents/Academics/Year 4/Semester 2/ECE
366 66100/hw09/stereoInputs/imgR.jpg", imgR.copy(), corrsR)
367
368 # Normalize the correspondences
369 normalizedptsL, tMatL = normalizepoints(corrsL)
370 normalizedptsR, tMatR = normalizepoints(corrsR)
371
372 # calculate the fundamental matrix
373 F = calculatefundamentalmatrix(normalizedptsL, normalizedptsR)
374 denormalizedF = tMatR.T @ F @ tMatL
375 denormalizedF = denormalizedF / denormalizedF[2, 2] # Make it homogeneous
376
377 # Find the epipoles and the camera projection matrices
378 eLnorm, eRnorm, Exnorm = findepipoles(denormalizedF)
379 PLnorm, PRnorm = computeprojectionmatrices(denormalizedF, eRnorm,
380 Exnorm)
381
382 # Using nonlinear least squares minimization find optimum F
383 f = np.ravel(denormalizedF)
384 Foptimized = leastsquares(costfunction, f, args=[corrsL, corrsR], method=
385 "lm").x
386 Foptimized = conditionF(np.reshape(Foptimized, (3, 3))) # Make the
387 fundamental matrix have a rank of 2
388 Foptimized = Foptimized / Foptimized[2, 2] # Make it homogeneous
389
390 # Get the optimized epipolar lines and camera projection matrices
391 eL, eR, Ex = findepipoles(Foptimized)
392 PL, PR = computeprojectionmatrices(Foptimized, eR, Ex)

```

```

388 # ======IMAGE RECTIFICATION
389 # ======
390 HL, HR = computehomographyrectification(imgL, corrsL, corrsR, eL, eR)
391 rectifiedleftimage = transform(imgL, HL)
392 cv2.imwrite(r"/Users/nikitaravi/Documents/Academics/Year 4/Semester 2/ECE
393   66100/hw09/rectified/imgL.jpg", rectifiedleftimage)
394
395 rectifiedrightimage = transform(imgR, HR)
396 cv2.imwrite(r"/Users/nikitaravi/Documents/Academics/Year 4/Semester 2/ECE
397   66100/hw09/rectified/imgR.jpg", rectifiedrightimage)
398
399 # ======INTEREST POINT DETECTION
400 # ======
401 # Detect Edges
402 edgesL = detectedges(rectifiedleftimage)
403 cv2.imwrite(r"/Users/nikitaravi/Documents/Academics/Year 4/Semester 2/ECE
404   66100/hw09/edges/imgL.jpg", edgesL)
405
406 edgesR = detectedges(rectifiedrightimage)
407 cv2.imwrite(r"/Users/nikitaravi/Documents/Academics/Year 4/Semester 2/ECE
408   66100/hw09/edges/imgR.jpg", edgesR)
409
410 # Detect Correspondences
411 corrsDetected = detectcorrespondences(edgesL, edgesR, maxsearcharea=15)
412 corrsFiltered = filtercorrespondences(rectifiedleftimage,
413   rectifiedrightimage, corrsDetected, kernelsize=15, maximumcorrs=500,
414   mode="ssd")
415 drawcorrespondences(rectifiedleftimage, rectifiedrightimage,
416   rectifiedleftimage.shape[1], corrsFiltered, path=r"/Users/nikitaravi/
417   Documents/Academics/Year 4/Semester 2/ECE 66100/hw09/correspondences/corr.
418   jpg")
419
420 # ======PROJECTIVE RECONSTRUCTION
421 # ======
422 rectifiedCorrsL = transformpts(corrsL, HL)
423 rectifiedCorrsR = transformpts(corrsR, HR)
424 rectifiedcoordinates = [[[xL, yL], [xR, yR]] for (xL, yL), (xR, yR) in zip(
425   rectifiedCorrsL, rectifiedCorrsR)]
426
427 drawcircles(r"/Users/nikitaravi/Documents/Academics/Year 4/Semester 2/ECE
428   66100/hw09/rectified/imgLcircles.jpg", rectifiedleftimage.copy(),
429   rectifiedCorrsL)
430 drawcircles(r"/Users/nikitaravi/Documents/Academics/Year 4/Semester 2/ECE
431   66100/hw09/rectified/imgRcircles.jpg", rectifiedrightimage.copy(),
432   rectifiedCorrsR)
433
434 unfilteredworldcoords = projectivereconstruction(rectifiedcoordinates,
435   PL, PR)
436 filteredworldcoords = projectivereconstruction(corrsFiltered, PL, PR)
437
438 # ======3D RECONSTRUCTION
439 # ======

```

```

423 fig = plt.figure()
424 ax = fig.add_subplot(111, projection="3d")
425 ax.scatter(unfilteredworldcoords[:, 0], unfilteredworldcoords[:, 1],
426     unfilteredworldcoords[:, 2])
427 ax.scatter(filteredworldcoords[:, 0], filteredworldcoords[:, 1],
428     filteredworldcoords[:, 2])
429
430 pairoflines = [[0, 1],
431     [0, 3],
432     [0, 6],
433     [1, 2],
434     [2, 3],
435     [2, 5],
436     [3, 4],
437     [4, 5],
438     [4, 6]] # Points that should form a line
439
440 for pair in pairoflines:
441     ax.plot([unfilteredworldcoords[pair[0]][0], unfilteredworldcoords[pair[1]][0]],
442             [unfilteredworldcoords[pair[0]][1], unfilteredworldcoords[pair[1]][1]],
443             [unfilteredworldcoords[pair[0]][2], unfilteredworldcoords[pair[1]][2]])
444
445 plt.show()
446
447 def getdmax(gtDisp, mul):
448     gtDisp = cv2.cvtColor(gtDisp, cv2.COLORBGR2GRAY)
449     gtDisp = gtDisp.astype(np.float32) / (16 * mul) # divide by mul because
450         disparity scale is 4
451     gtDisp = gtDisp.astype(np.uint8)
452
453     dmax = np.max(gtDisp)
454     return gtDisp, dmax
455
456 def censustransform(imgL, imgR, M, dmax):
457     height, width = imgL.shape[:2]
458
459     grayL = cv2.cvtColor(imgL, cv2.COLORBGR2GRAY) if len(imgL.shape) == 3 else
460         imgL
461     grayR = cv2.cvtColor(imgR, cv2.COLORBGR2GRAY) if len(imgR.shape) == 3 else
462         imgR
463
464     halfM = M // 2
465     bordersize = dmax + halfM
466     disparitymap = np.zeros((height, width), dtype=np.uint8)
467
468     for rowL in range(bordersize, height - bordersize):
469         print(f"Row: {rowL} / {height - bordersize}")
470         for colL in range(width - bordersize - 1, bordersize - 1, -1):
471             cost = []
472             leftWindow = grayL[rowL - halfM: rowL + halfM + 1,
473                             colL - halfM: colL + halfM + 1]

```

```

469     binaryleftwindow = np.ravel(np.where(leftWindow & grayL[rowL, colL],
470     1, 0))
471     for d in range(dmax + 1):
472         rowR, colR = rowL, colL - d
473         rightWindow = grayR[rowR - halfM: rowR + halfM + 1,
474             colR - halfM: colR + halfM + 1]
475
476         binaryrightwindow = np.ravel(np.where(rightWindow & grayR[rowL, colR],
477             1, 0))
478         cost.append(np.sum(np.bitwiseor(binaryleftwindow,
479             binaryrightwindow)))
480
481     disparitymap[rowL, colL] = np.argmin(cost)
482
483 return disparitymap.astype(np.uint8)

484 def computedisparityaccuracy(gtDisp, dispMap, dmax, M, delta=2):
485     halfM = M // 2
486     bordersize = halfM + dmax
487     dispMap = dispMap[bordersize:dispMap.shape[0]-bordersize, bordersize:
488         dispMap.shape[1]-bordersize]
489     gtDisp = gtDisp[bordersize:gtDisp.shape[0]-bordersize, bordersize:gtDisp.
490         shape[1]-bordersize]
491
492     error = np.abs(dispMap.astype(np.uint16) - gtDisp.astype(np.uint16)).astype(
493         np.uint8)
494     validpixels = cv2.countNonZero(gtDisp) # Non-black pixels are valid pixels
495     # 74952
496     print("The number of valid pixels: ", validpixels)
497
498     accuracy = np.sum(error != delta)
499     percentageofaccuracy = accuracy / validpixels
500
501     errormask = np.where(error != delta, 255, 0)
502     return percentageofaccuracy, errormask

503 def densestereomatching(M=10):
504     imgL = cv2.imread(r"/Users/nikitaravi/Documents/Academics/Year 4/Semester 2/
505         ECE 66100/hw09/Task3Images/im2.ppm")
506     dispL = cv2.imread(r"/Users/nikitaravi/Documents/Academics/Year 4/Semester
507         2/ECE 66100/hw09/Task3Images/disp2.pgm")
508
509     imgR = cv2.imread(r"/Users/nikitaravi/Documents/Academics/Year 4/Semester 2/
510         ECE 66100/hw09/Task3Images/im6.ppm")
511     dispR = cv2.imread(r"/Users/nikitaravi/Documents/Academics/Year 4/Semester
512         2/ECE 66100/hw09/Task3Images/disp6.pgm")
513
514     gtDisp, dmax = getdmax(dispL, mul=0.25) # dmax = 13 * mul
515     print("dmax: ", dmax)
516
517     disparitymap = censustransform(imgL, imgR, M, dmax)

```

```

511 savedisparitymap = cv2.normalize(disparitymap, dst=None, alpha=0, beta
512     =255, normtype = cv2.NORMMINMAX).astype(np.uint8)
513 cv2.imwrite(r"/Users/nikitaravi/Documents/Academics/Year 4/Semester 2/ECE
514     66100/hw09/disparityMap/" + str(M) + "size.jpg", savedisparitymap)
515
516 percentageofaccuracy, errormask = computedisparityaccuracy(gtDisp,
517     disparitymap, dmax, M)
518 print("Percentage of Accuracy: ", percentageofaccuracy)
519 cv2.imwrite("/Users/nikitaravi/Documents/Academics/Year 4/Semester 2/ECE
520     66100/hw09/errormask/errormask" + str(M) + ".jpg", errormask)
521
522
523 if name == " main ":
524     # TASK 1
525     projectivestereoreconstruction()
526
527     # TASK 2
528     densestereomatching(M=30)
529     densestereomatching(M=45)
530     densestereomatching(M=60)

```

Listing 1: The Source Code

References

- [1] R. Hartley, A. Zisserman, *Multiple View Geometry in Computer Vision*, Cambridge, 2003.
- [2] A. Kak, “A Loop and Zhang Reader for Stereo Rectification”,
<https://engineering.purdue.edu/kak/Tutorials/StereoRectification.pdf> 2022.